

Deep Learning Assignment

Swapnil Anand(21285)

March 27, 2024

Question 1

In logistic regression, the cross-entropy loss function is typically preferred over mean squared error (MSE) for several reasons, including its ability to provide a clearer and more optimal solution.

Nature of Logistic Regression:

Logistic regression is commonly used for binary classification problems, where the target variable has two possible outcomes (e.g., 0 or 1). The logistic function (also known as the sigmoid function) is used to model the probability that a given input belongs to one of the classes.

Interpretability:

Cross-entropy loss directly relates to the likelihood of the predicted probabilities matching the actual labels. It's intuitive to interpret: the lower the cross-entropy loss, the better the model's predictions align with the true labels. In contrast, MSE doesn't directly relate to the probabilities. It's more suited for regression tasks where the output is continuous and measuring the squared difference between predicted and actual values is meaningful.

Gradient Descent Optimization:

Cross-entropy loss leads to smoother gradients during optimization compared to MSE, especially when used in conjunction with the logistic activation function. In logistic regression, optimizing cross-entropy loss through techniques like gradient descent tends to converge faster and more reliably to the global optimum due to its convex nature.

Single Optimal Solution:

Cross-entropy loss guarantees a single best solution due to its convexity. There's no ambiguity in the optimization process, and the model converges to a unique set of parameters that minimize the loss function. With MSE, especially in logistic regression where it's not the ideal loss function, there might be multiple local minima, leading to potential confusion and suboptimal solutions.

Question 2

In a binary classification task with a deep neural network equipped with linear activation functions, the loss function that guarantees a convex optimization problem is (b) Mean Squared Error (MSE).

Mean Square Error

In binary classification, let y denote the target label (either 0 or 1) and \hat{y} denote the predicted output of the model. With linear activation functions, the output of the model is a linear combination of inputs and weights, i.e

$$\hat{y} = WX + b$$

As the output of the model is a linear function of the inputs, the MSE loss function becomes a quadratic function of the parameters (weights and biases). Quadratic functions are convex, meaning they have a single global minimum. Thus, using MSE as the loss function guarantees convexity in the optimization problem.

Cross Entropy

While Cross Entropy loss is commonly used for binary classification, it's not guaranteed to result in a convex optimization problem when combined with linear activation functions. The nonlinearity introduced by the logarithm function makes the loss function non-convex, and the optimization problem may have multiple local minima.

Question3

We define a simple feedforward neural network (FNN) with three dense layers.

The input images are preprocessed using torchvision transforms, which convert images to tensors and normalize them. We use the MNIST dataset for training and testing. The FNN model is trained using stochastic gradient descent (SGD) with cross-entropy loss. We evaluate the trained model on the test set and calculate its accuracy. Hyperparameters such as learning rate, number of hidden layers, number of neurons per layer, and activation functions (ReLU in this case) can be tuned to optimize the performance of the model. Techniques like grid search or random search can be employed for hyperparameter tuning. Additionally, techniques like dropout or batch normalization can be used to prevent overfitting and improve generalization.

```

+ Code + Text
print("Accuracy on the test set: {}%".format(100 * correct / total))

Downloading https://www.jocan.com/cdn/mnist/train-images-idx3-ubyte.gz
Downloading https://www.jocan.com/cdn/mnist/train-labels-idx1-ubyte.gz
Extracting ./data/MNIST/train-images-idx3-ubyte.gz to ./data/MNIST/train
Extracting ./data/MNIST/train-labels-idx1-ubyte.gz to ./data/MNIST/train

Downloading https://www.jocan.com/cdn/mnist/10k-images-idx3-ubyte.gz
Downloading https://www.jocan.com/cdn/mnist/10k-labels-idx1-ubyte.gz
Extracting ./data/MNIST/10k-images-idx3-ubyte.gz to ./data/MNIST/10k
Extracting ./data/MNIST/10k-labels-idx1-ubyte.gz to ./data/MNIST/10k

Epoch 1, Batch 100, Loss: 2.207
Epoch 1, Batch 200, Loss: 1.879
Epoch 1, Batch 300, Loss: 1.388
Epoch 1, Batch 400, Loss: 0.967
Epoch 1, Batch 500, Loss: 0.751
Epoch 1, Batch 600, Loss: 0.685
Epoch 1, Batch 700, Loss: 0.550
Epoch 1, Batch 800, Loss: 0.497
Epoch 1, Batch 900, Loss: 0.460
Epoch 2, Batch 100, Loss: 0.436

```

Figure 1: Question 3 image

```

+ Code + Text
epoch 3, Batch 100, Loss: 0.140
epoch 3, Batch 200, Loss: 0.110
epoch 3, Batch 300, Loss: 0.122
epoch 3, Batch 400, Loss: 0.112
epoch 3, Batch 500, Loss: 0.125
epoch 3, Batch 600, Loss: 0.107
epoch 3, Batch 700, Loss: 0.200
epoch 3, Batch 800, Loss: 0.207
epoch 3, Batch 900, Loss: 0.204
epoch 4, Batch 100, Loss: 0.280
epoch 4, Batch 200, Loss: 0.284
epoch 4, Batch 300, Loss: 0.280
epoch 4, Batch 400, Loss: 0.281
epoch 4, Batch 500, Loss: 0.284
epoch 4, Batch 600, Loss: 0.281
epoch 4, Batch 700, Loss: 0.283
epoch 4, Batch 800, Loss: 0.291
epoch 4, Batch 900, Loss: 0.291
epoch 5, Batch 100, Loss: 0.270
epoch 5, Batch 200, Loss: 0.276
epoch 5, Batch 300, Loss: 0.277
epoch 5, Batch 400, Loss: 0.257
epoch 5, Batch 500, Loss: 0.251
epoch 5, Batch 600, Loss: 0.251
epoch 5, Batch 700, Loss: 0.270
epoch 5, Batch 800, Loss: 0.289
epoch 5, Batch 900, Loss: 0.256
Finished Training
Accuracy on the test set: 92.7%

```

Figure 2: Question 3 image 2

Question4

For Lenet we trained for 100 mini batches for 5 epochs. After training the model we achieved the accuracy of 87.53. The other models were very time consuming.